

VibeValue: Language Model for Enhanced Financial Sentiment Insights

Diptadeep Sinha
VIT Bhopal University
Bhopal, India
khdiptadeep@gmail.com

Preet Dadga
VIT Bhopal University
Bhopal, India
dadgapreet9426@gmail.com

Prithvi Narayan Singh
VIT Bhopal University
Bhopal, India
prithviofficial99@gmail.com

Parth Raj
VIT Bhopal University
Bhopal, India
rajparth8776@gmail.com

Pragjyotish Lahan
VIT Bhopal University
Bhopal, India
pragjyotishlahan@gmail.com

Abstract—Pretrained language models with contextual understanding, have achieved notable advancements in diverse NLP tasks through training on extensive amounts of unlabeled text. The financial sector generates substantial volumes of financial communication text. Despite this, there is currently a lack of pretrained language models specifically tailored for finance. In this study, we aim to fill this gap by using VibeValue a product after fine-tuning a pre-trained sentiment analysis model on a finance domain-specific data and optimizing its output to deliver more accurate results to a specified dataset.

Index Terms—

I. INTRODUCTION

Financial markets are dynamic ecosystems where billions of transactions are initiated and completed daily, making them a focal point of global economic activity. Prices in an open market reflects all of the available information regarding assets exchanged in an economy. [12] The ability to comprehend and respond to the intricate interplay of factors influencing these markets is a fundamental challenge for investors, traders, and financial institutions. Analysis of financial texts, be it news, analyst reports or official company announcements is a possible source of new information with an unprecedented amount of such text being created every day, manually analyzing these and deriving actionable insights from them is too big of a task for any single entity. Hence, automated sentiment or polarity analysis of texts produced by financial actors using natural language processing (NLP) methods has gained popularity during the last decade. [5] Sentiment analysis is the task of extracting sentiments or opinions of people from written language. [9]The two paradigms of recent effort are:

- 1) Machine learning methods with features extracted from text with "word counting" [1] [2] [17] [19].
- 2) Deep learning methods, where text is represented by a sequence of embeddings [3] [16] [21]. The former suffers from inability to represent the semantic information that results from a particular sequence of words, while

the latter is often deemed as too "data-hungry" as it learns a much higher number of parameters [2].

VibeValue leverages the architectural foundation of BERT [4] and distinguishes itself by being pre-trained on an extensive corpus of financial communication data, equipping it with a deep understanding of financial jargon, market-specific phrases, and nuances that generic language models may overlook. This domain-specific fine-tuning of our model positions it as a pivotal asset for analyzing financial sentiment. BERT, the foundation upon which VibeValue is built, stands as one of the most influential breakthroughs in NLP. It embodies the essence of contextual modeling by considering the entire sentence and context of a word when predicting its meaning. This approach contrasts with traditional methods that rely on fixed-word embeddings or bag-of-words representations, which lack the context-awareness necessary for capturing nuanced sentiment in financial texts. The choice to utilize VibeValue for market sentiment prediction over traditional methods is motivated by several compelling reasons. Traditional approaches such as Using carefully crafted financial sentiment lexicons such as Loughran and McDonald (2011) [10] may seem a solution because they incorporate existing financial knowledge into textual analysis. However, they are based on "word counting" methods, which come short in analyzing deeper semantic meaning of a given text. These methods struggle to capture the subtleties of financial language and the context-dependent nature of sentiment expression. In contrast, our model excels in recognizing the intricate sentiment patterns within financial texts, thanks to its pre-training on a vast financial corpus. Its contextual understanding allows it to discern the nuances of market-related language, including industry-specific terminology, slang, and metaphors. This contextual awareness is a key advantage over traditional sentiment analysis tools, as it enables VibeValue to provide more accurate and insightful sentiment assessments, the deep learning architecture employed makes it adaptable to evolving market dynamics and emerging linguistic trends,

traditional methods often require manual adjustments and rule-based updates, making them less agile in responding to changing market conditions and evolving language usage. The journey of "VibeValue" takes us into the heart of NLP and its application in the high stakes world of financial markets. Through advanced sentiment analysis, we seek to empower investors, financial analysts, and institutions with a tool that aids in making informed, data-driven decisions in an increasingly complex and interconnected global economy. In the subsequent sections of this project, we delve into the methodology, results, and implications of our fine-tuned model as the sentiment analysis of financial text data using other state of the art models present challenges in achieving high accuracy, precision, recall, and F1 score for this reason fine-tuning such base models is required to address these limitations. To illustrate that our proposed model can enhance sentiment analysis by improving the accuracy of sentiment classification, ensuring that the model comprehensively understands and interprets the nuanced language inherent to the financial domain and contribute to a more competitive, insightful, and robust financial landscape.

II. RELATED WORK

A. LSTM

Long Short-Term Memory (LSTM) networks represent a specialized type of recurrent neural network designed to retain long-term dependencies within sequences. They achieve this by incorporating mechanisms such as "forget" gate which decides what information from the cell state should be thrown away or kept. It takes as input the previous hidden state h_{t-1} and the current input x_t , and outputs a value between 0 and 1 for each number in the cell state C_{t-1} . A value of 1 means "keep this" while a value of 0 means "get rid of this." Mathematically:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

The input gate decides what new information should be stored in the cell state. It consists of two parts: a sigmoid layer that decides which values to update, and a tanh layer that creates a vector of new candidate values to add to the state.

Mathematically:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\beta = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

and "update" gates which combines the decisions from the forget and input gates to update the old cell state C_{t-1} into the new cell state C_t .

Mathematically:

$$C_t = f_t \cdot C_{t-1} + i_t \cdot \beta$$

LSTMs have emerged as foundational architectures for modeling a diverse array of sequential data, ranging from stock price movements to natural language processing tasks.

When applying LSTM-based models to natural language tasks, such as text processing, an essential consideration involves determining the initial representation for individual tokens within the sequence. A prevalent approach entails leveraging pre-trained weights for these initial token representations. One widely recognized pre-training method is GLoVe [14] (Global Vectors for Word Representation), which focuses on generating word embeddings through an unsupervised learning process involving a log-bilinear regression model trained on a word-word co-occurrence matrix derived from extensive text corpora. While GLoVe effectively captures semantic relationships between words within a vector space, in which the similarity between word vectors can be calculated using cosine similarity or other distance metrics. Mathematically, cosine similarity between word vectors w_i and w_j is given by:

$$\text{similarity}(w_i, w_j) = \frac{w_i \cdot w_j}{\|w_i\| \cdot \|w_j\|}$$

it lacks the capability to contextualize these representations within the specific sequences in which the words appear, underscoring the need for more context-aware modeling techniques.

B. ELMo

ELMo embeddings [15] provide contextualized word representations, wherein the context of surrounding words influences the embedding of each word. At the core of ELMo lies a bidirectional language model, incorporating multiple layers of LSTM units. The primary objective of this language model is to capture the probability distribution across sequences of tokens within a specified vocabulary. the word embeddings are computed as:

$$ELMo(c_t) = E.(r_{word} \cdot \vec{h}_t + h_t^{\leftarrow})$$

where E is a linear transformation, and r_{word} is a vector of learned parameters.

Specifically, ELMo evaluates the likelihood of a token based on both preceding and subsequent tokens in the sequence. Furthermore, the model dynamically adjusts the weights assigned to various representations across different LSTM layers to generate a singular, context-enhanced vector for each token. Once these contextually enriched representations are derived, they can serve as valuable initializations for various downstream natural language processing tasks.

ELMo captures contextual information by aggregating the entire sequence of hidden states $\vec{h}_1, \vec{h}_2, \dots, \vec{h}_T$ and $h_1^{\leftarrow}, h_2^{\leftarrow}, \dots, h_T^{\leftarrow}$ using layer-wise summation. Mathematically, the contextual embedding ELMo(x) for the entire input sequence x is computed as:

$$ELMo(x) = E.(\sum_{t=1}^T r_{context}^t \cdot (\vec{h}_t) + h_t^{\leftarrow})$$

where $r_{context}^t$ is a vector of learned parameters for each layer t.

C. ULMFit

ULMFit is a transfer learning model for down-stream NLP tasks, that make use of language model pre-training [6]. Unlike ELMo, with ULMFit, the whole language model is fine-tuned together with the task-specific layers. The underlying language model used in ULMFit is AWD-LSTM, which uses sophisticated dropout tuning strategies to better regularize its LSTM model [4]. For classification using ULMFit two linear layers are added to the pre-trained AWD-LSTM, first of which takes the pooled last hidden states as input. Mathematically, given a sequence of input tokens $x_1, x_2, \dots, x_T$ and their corresponding target tokens $y_1, y_2, \dots, y_T$, the objective is to minimize the cross-entropy loss:

$$Loss = -\sum_{t=1}^T \log P(y_t | x_1, x_2, \dots, x_{t-1})$$

ULMFit comes with novel training strategies for further pretraining the language model on domain-specific corpus and finetuning on the down-stream task.

ULMFit uses slanted triangular learning rates, where the learning rate is increased linearly for a certain portion of training epochs and then decreased exponentially. This technique helps in faster convergence and better generalization. Mathematically, the learning rate schedule is defined as:

$$\eta(t) = \frac{\eta_{max} - \eta_{min}}{2 \cdot sl} \cdot \min(step, sl - step) + \eta_{min}$$

where η_{max} is the maximum learning rate, η_{min} is the minimum learning rate, and sl is the percentage of steps to increase the learning rate linearly.

D. Transformers

The Transformer represents an attention-driven architecture designed for processing sequential data, offering a compelling alternative to recurrent neural networks [7]. Initially conceived as a sequence-to-sequence model, the Transformer incorporates both encoder and decoder components; however, our discussion will predominantly center on the encoder component, noting its analogous structure with the decoder. The encoder module comprises multiple identical Transformer layers, each encompassing a multi-headed self-attention mechanism which allows the model to weigh the importance of different input tokens when producing an output token. It computes the attention scores by taking the dot product of the query, key, and value vectors, followed by scaling and softmax operation. The output is then computed as a weighted sum of the values, where the weights are determined by the attention scores.

Mathematically:

- Attention Score:

$$Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_k}})$$

- Output:

$$Attention(Q, K, V) \cdot V$$

paired with a fully connected feed-forward network which applies two linear transformations separated by a non-linear activation function, usually ReLU.

Mathematically:

$$FFN(x) = ReLU(xW_1 + b_1)W_2 + b_2$$

Within a self-attention layer, the model learns three distinct mappings from embeddings, namely key, query, and value. Utilizing the key associated with each token and the query vectors derived from all tokens, the architecture computes a similarity score via dot product operations. These computed scores subsequently dictate the weighting applied to the value vectors, culminating in a refined representation for each token. Through the integration of multi-headed self-attention across layers, the Transformer allows for a multifaceted evaluation of the sequence, enabling diverse perspectives to inform the representation. Finally, these processed vectors traverse shared parameterized fully connected networks, further enhancing the model's capacity to capture intricate sequential relationships. As it was argued by Vaswani 2017 [18], Transformer architecture has several advantages over the RNN-based approaches. Because of RNNs' sequential nature, they are much harder to parallelize on GPUs and too many steps between far away elements in a sequence make it hard for information to persist

E. BERT

BERT [1] operates as a groundbreaking language model, featuring a series of stacked Transformer encoders and diverging from traditional language modeling methods like ELMo and AWD-LSTM. Instead of predicting subsequent words from preceding ones, BERT adopts a "masking" approach by randomly concealing approximately 15% of tokens within sequences, aiming to predict these obscured tokens using a softmax layer. Concurrently, BERT is trained on a "next sentence prediction" task to determine the logical continuity between pairs of sentences. For input representation, BERT incorporates token and positional embeddings, supplemented by special tokens—[CLS] and [SEP]—appended to sequence beginnings and ends, respectively, with the [CLS] token serving a crucial role in various classification tasks. BERT is available in two versions: BERT-base, featuring 12 encoder layers, a hidden size of 768, 12 multi-head attention heads, and 110 million parameters; and BERT-large, with 24 encoder layers, a hidden size of 1024, 16 attention heads, and 340 million parameters. Both variants are trained on extensive datasets from BookCorpus [22] and English Wikipedia, encompassing over 3.5 billion words.

III. PROPOSED MODEL

On identifying the need for refined and specialized sentiment analysis tool for financial text, our research introduces a novel model named VibeValue to address the nuances of financial domain specific vocabulary and complex sentiments on the market. Built upon the foundational architecture of FinBERT [2] [20], VibeValue represents a specialized deep learning model, fine-tuned to excel in the domain of financial

communication. The core idea behind these models is that by training language models on very large corpora and then initializing down-stream models with the weights learned from the language modeling task, a much better performance can be achieved. The initialized layers can range from the single word embedding layer [6] to the whole model [15]. This approach should, in theory, be an answer to the scarcity of labeled data problems. Language models don't require any labels, since the task is predicting the next word. They can learn how to represent the semantic information. That leaves the fine-tuning on labeled data only the task of learning how to use this semantic information to predict the labels. Our model undergoes domain-specific further pretraining on an extensive financial corpus, this is to employ one particular component of transfer learning methods, which is the ability to further pretrain the language model, enabling it to capture intricate language patterns and sentiments niche to the financial domain. The model introduces adaptations in the token embeddings, attention mechanisms, and output layers to optimize its performance for classifying sentiment in financial texts. Furthermore, VibeValue incorporates a task-specific fine-tuning process using labeled financial datasets, allowing it to adapt to the nuances of diverse financial contexts. The architecture features specialized layers to handle financial jargon and entities, enhancing its sensitivity to the subtleties of sentiment expression in various financial communication channels. The model's innovative design extends beyond traditional sentiment analysis, offering a nuanced understanding of financial language and sentiments, making it a valuable asset for decision-makers navigating the complexities of financial information. Through systematic evaluation and comparison with baseline models, VibeValue demonstrates its effectiveness in achieving superior accuracy, precision, and F1 scores, establishing its position as a powerful and versatile tool for sentiment analysis in the financial sector.

IV. MODULE DESIGN AND IMPLEMENTATION

The implementation of our project VibeValue a sophisticated sentiment analysis model tailored for the financial domain employs a pre-trained BERT model on financial datasets comprising of a huge financial corpus to introduce domain specific vocabulary of the finance sector and using it as our base model on which we employ our fine-tuning strategies by setting up and optimizing the hyper parameters and selecting specific combinations of datasets to employ transfer learning. The process of fine-tuning our model on different annotated datasets for fine-tuning the classification layer for sentiment analysis viz. Financial PhraseBank [13], a combined dataset of FiQA [11] and Financial PhraseBank. Our model is engineered to capture the nuances and context of financial language, thereby enhancing its ability to accurately assess market sentiment. In addition to fine-tuning, we optimize hyperparameters and apply techniques to handle imbalanced datasets, ensuring robust performance. This model forms the core of VibeValue, providing a powerful tool for real-time sentiment analysis in financial markets, empowering users to

make informed decisions in an environment where sentiment holds the key to success.

V. DATASETS

A. Financial PhraseBank

The main sentiment analysis dataset used to train our model is Financial PhraseBank from Malo et al. 2014 [13]. This release of the financial phrase bank covers a collection of 4840 sentences. The selected collection of phrases was annotated by 16 people with adequate background knowledge on financial markets. Three of the annotators were researchers and the remaining 13 annotators were master's students at Aalto University School of Business with majors primarily in finance, accounting, and economics. Given the large number of overlapping annotations (5 to 8 annotations per sentence), there are several ways to define a majority vote-based gold standard. To provide an objective comparison, we have formed 4 alternative reference datasets based on the strength of majority agreement:

- sentences with 100% agreement
- sentences with more than 75% agreement
- sentences with more than 66% agreement
- sentences with more than 50% agreement

B. A combined Dataset of FiQA and Financial PhraseBank

FiQA [11] is a dataset that was created for WWW '18 conference financial opinion mining and question answering challenges. The Dataset that we have used has 5843 annotated sentences which is used to train the model.

VI. MODEL ARCHITECTURE

A. Data Collection

The finetuning process employs the Financial PhraseBank dataset by Malo et al. (2014) [1] for sentiment analysis. Comprising 4845 English sentences randomly drawn from financial news on LexisNexis, the dataset underwent annotation by 16 finance and business experts, who labeled sentences based on perceived impacts on company stock prices. Information on annotator agreement levels is included. Test and validation sets were created, with 20% of sentences allocated to each. The training set consists of 3101 examples. Our model VibeValue was employed and further trained on this dataset.

B. Layers and components relevant to sentiment analysis

VibeValue is built upon BERT architecture. BERT has 12 Transformer encoder layers. Each encoder layer has components that play a crucial role in understanding and processing the input text, which is especially relevant for sentiment analysis in the financial domain. The encoder consists of multiple identical Transformer layers. Each layer has a multi-headed self-attention layer and a fully connected feed-forward network it also processes the input sequence in a self-attention mechanism, allowing the model to weigh the importance of each word in relation to the others. The multiple layers of the transformer encoder enable the model to capture intricate contextual information in financial text as financial

language often involves complex structures, and the layer-wise processing helps in understanding the nuanced relationships between words. The input embedding layer is responsible for converting words into numerical vectors (word embeddings). BERT utilizes WordPiece tokenization, breaking down words into subword units for a more granular representation; this layer helps in representing the financial text in a format that the subsequent layers of the model can comprehend. The context-aware embeddings capture the relationships between words, which is crucial for sentiment analysis. Self-attention layer allows each word in a sequence to focus on different parts of the input sequence, enabling the model to consider the context in both directions (bidirectional context). The self-attention mechanism allows VibeValue to weigh the importance of different words based on their relevance to sentiment expression, capturing dependencies and nuances in the financial text. The final layers of the model produce output representations for downstream tasks, such as sentiment analysis. These representations can be used to classify the sentiment of the input text. The output layer is where the model generates predictions regarding the sentiment of the financial text. Depending on the specific implementation, this could involve a classification into positive, negative, or neutral sentiments.

C. Training setup and parameters

For the finetuning process after initializing the base model, we configured the hyperparameters. The first training iteration utilized a mini batch size of 32, the training process ran for 20 epochs, and employed a learning rate of $2e-5$ on Financial PhraseBank dataset which was divided into training (80%) and testing (20%) subsets while maintaining reproducibility through a fixed random seed. We documented the model's performance scores and saved it as "vibevalue1" for further training on our second dataset, FiQA+FB. The hyperparameters remained consistent except for a reduced batch size of 16, and this model was saved as "vibevalue2." We then looped back and trained "vibevalue2" on our first dataset, documenting its performance scores as "vibevalue3." This process demonstrated the model's adaptability to new data sources, highlighting the effectiveness of transfer learning in financial sentiment analysis. In subsequent comparisons with earlier iterations, "vibevalue3" is denoted as "VibeValue".

D. Evaluation metrics used for performance assessment

In this subsection, we detail the evaluation metrics employed to assess the performance of VibeValue, a sentiment analysis model specifically fine-tuned for the financial domain. The choice of appropriate metrics is critical for comprehensively evaluating the model's ability to discern and classify sentiments within financial texts.

1) *Accuracy*: Accuracy serves as a fundamental metric, providing a holistic view of the model's overall correctness in sentiment classification. It is calculated as the ratio of correctly predicted instances to the total number of instances in the dataset. A high accuracy score indicates the model's




Fig. 1. Figure: Accuracy Variation of VibeValue with increasing epochs

proficiency in making correct predictions across all sentiment classes.

2) *Precision, Recall, and F1 Score*: Precision, recall, and F1 score are particularly important for sentiment analysis where the balance between positive, negative, and neutral sentiments can vary. Precision measures the accuracy of positive predictions, recall quantifies the model's ability to capture all positive instances, and F1 score balances the trade-off between precision and recall. These metrics are especially relevant for applications where misclassifications in either direction carry significant implications, such as financial decision-making.

VII. EXPERIMENTAL RESULTS

A. Fine-tuned Model Performance

The analysis of accuracy variation with increasing epochs for VibeValue reveals a compelling trajectory in the model's learning process. As epochs progress, there is a discernible pattern of accuracy improvement, indicating the model's ability to refine its understanding of financial sentiment over time. This upward trend suggests that VibeValue benefits from continued exposure to the training data, refining its internal representations and adapting to the intricacies of financial language. The observed accuracy variation underscores the importance of an extended training duration for VibeValue, allowing it to capture nuanced patterns and subtleties in sentiment across diverse financial contexts. Fine-tuning and iterative learning through increasing epochs contribute to the model's overall proficiency in financial sentiment analysis, highlighting the dynamic nature of its learning process.

The examination of training and validation loss dynamics in VibeValue over increasing epochs provides valuable insights into the model's learning behavior. As the number of epochs increases, the training loss demonstrates a decreasing trend, indicating that the model is effectively minimizing errors on the training dataset. This reduction in training loss signifies improved convergence and suggests that the model is successfully




Fig. 2. Loss variation of VibeValue with increasing epochs

TABLE I
COMPARISON OF VIBEVALUE (VIBEVALUE3) WITH VARIOUS NLP
MODELS ON COMPLETE DATASET

Sr No.	Model	All Data		
		Loss	Accuracy	F1 Score
1	LSTM	0.81	0.71	0.64
2	LSTM with ELMo	0.72	0.75	0.7
3	ULMFit	0.41	0.83	0.79
4	LPS	-	0.71	0.71
5	HSC	-	0.71	0.76
6	FinSSLX	-	-	-
7	FinBert	0.37	0.86	0.84
8	VibeValue	0.2539	0.9019	0.888

learning to generalize patterns from the training data. Simultaneously, monitoring the validation loss is crucial for assessing the model's performance on unseen data. The validation loss initially decreases, signifying improved generalization, but it may eventually stabilize or even increase, indicating potential overfitting. The interplay between training and validation loss across epochs provides a nuanced perspective on VibeValue's learning process, aiding in the fine-tuning and optimization of the model for robust financial sentiment analysis.

The results of VibeValue , the baseline methods and state-of-the-art on Financial PhraseBank dataset classification task

TABLE II
COMPARISON OF VIBEVALUE (VIBEVALUE3) WITH VARIOUS NLP
MODELS ON DATA WITH 100% AGREEMENT

Sr No.	Model	Data with 100% agreement		
		Loss	Accuracy	F1 Score
1	LSTM	0.57	0.81	0.74
2	LSTM with ELMo	0.50	0.84	0.77
3	ULMFit	0.20	0.93	0.91
4	LPS	-	0.79	0.80
5	HSC	-	0.83	0.86
6	FinSSLX	-	0.91	0.85
7	FinBert	0.13	0.97	0.95
8	VibeValue	0.063	0.9801	0.973




Fig. 3. Figure: Comparison of VibeValue accuracy with state-of-the-art NLP models

can be seen on table. We present the result on both the whole dataset and subset with 100% annotator agreement. The Table summarizes the results of VibeValue, baseline methods (LSTM and ULMFit), and state-of-the-art models (LPS, HSC, FinSSLX) on the entire Financial PhraseBank dataset. Notably, VibeValue demonstrates superior performance across all measured metrics compared to both the implemented methods (LSTM and ULMFit) and models reported in other studies[] (LPS [13], HSC [8], FinSSLX [11]). LSTM classifier with no language model information performs the worst. In terms of accuracy, it is close to LPS and HSC, (even better than LPS for examples with full agreement), however it produces a low F1-score. That is due to it performing much better in neutral class. LSTM classifier with ELMo embeddings improves upon LSTM with static embeddings in all of the measured metrics. It still suffers from low average F1-score due to poor performance in less represented labels. But it's performance is comparable with LPS and HSC, besting them in accuracy. So contextualized word embeddings produce close performance to machine learning based methods for dataset of this size. ULMFit significantly improves on all of the metrics and it doesn't suffer from model performing much better in some classes than the others. It also handily beats the machine learning based models LPS and HSC. This shows the effectiveness of language model pre-training. AWD-LSTM is a very large model and it would be expected to suffer from over-fitting with this small of a dataset. But due to language model pre-training and effective training strategies, it is able to overcome small data problem. ULMFit also outperforms FinSSLX, which has a text simplification step as well as pre-training of word embeddings on a large financial corpus with sentiment labels. VibeValue outperforms ULMFit, and consequently all of the other methods in all metrics.

The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In




Fig. 4. Figure: Comparison of VibeValue F1 score with state-of-the-art NLP models

extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset.

The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset.

performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset. The comparative analysis of VibeValue (vibevalue3) against various state-of-the-art Natural Language Processing (NLP) models unequivocally establishes VibeValue as a frontrunner, outperforming its counterparts across diverse benchmarks. In extensive evaluations, VibeValue consistently exhibits superior performance in tasks related to financial sentiment analysis. Notably, when compared with established NLP models, VibeValue consistently achieves higher accuracy, precision, recall, and F1-score metrics on the same dataset.

VIII. CONCLUSION

In conclusion, the findings of our study showcase the effectiveness of VibeValue in accurately classifying sentiment across diverse financial texts. The model's commendable performance, as reflected in precision, recall, and F1-score metrics, establishes its reliability and applicability as a specialized deep learning tool for financial sentiment analysis. This research contributes significantly to the field by not only demonstrating VibeValue's high accuracy but also enhancing the interpretability of other models. The comparative analysis framework developed serves as a benchmark for evaluating the relative strengths and weaknesses of various sentiment analysis approaches with state of the art models in the financial domain. Looking ahead, future work involves exploring multimodal analysis with additional data sources and types which can provide a holistic understanding of market sentiment, enhancing real-time analysis capabilities with parallel processing and efficient data pipelines, and further improving interpretability. Additionally, we propose extending the research into sentiment forecasting on how trends can evolve, conducting comparative studies against other deep learning models, traditional sentiment analysis methods, or domain-specific lexicon-based approaches, exploring transfer learning, and addressing advanced risk management strategies, behavioral finance integration, and

robustness against adversarial testing. These avenues of future work aim to advance the capabilities of sentiment analysis tools in navigating the complexities of financial markets and decision-making.

IX. REFERENCES

REFERENCES

- [1] Basant Agarwal and Namita Mittal. Machine learning approach for sentiment analysis. pages 21–45, 2016.
- [2] Dogu Araci. Finbert: Financial sentiment analysis with pre-trained language models. 8 2019.
- [3] Oscar Araque, Ignacio Corcuera-Platas, J. Fernando Sánchez-Rada, and Carlos A. Iglesias. Enhancing deep learning sentiment analysis with ensemble techniques in social applications. *Expert Systems with Applications*, 77:236–246, 7 2017.
- [4] Jacob Devlin, Ming Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *NAACL HLT 2019 - 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies - Proceedings of the Conference*, 1:4171–4186, 10 2018.
- [5] Li Guo, Feng Shi, and Jun Tu. Textual analysis and machine learning: Crack unstructured data in finance and accounting. *Journal of Finance and Data Science*, 2:153–170, 9 2016.
- [6] Jeremy Howard and Sebastian Ruder. Universal language model fine-tuning for text classification. pages 328–339.
- [7] Mathias Kraus and Stefan Feuerriegel. Decision support from financial disclosures with deep neural networks and transfer learning. *Decision Support Systems*, 104:38–48, 12 2017.
- [8] Srikumar Krishnamoorthy. Sentiment analysis of financial news articles using performance indicators. *Knowledge and Information Systems*, 56:373–394, 8 2018.
- [9] Bing Liu. Sentiment analysis and opinion mining. 2012.
- [10] Tim Loughran and Bill McDonald. When is a liability not a liability? textual analysis, dictionaries, and 10-ks. *Journal of Finance*, 66:35–65, 2 2011.
- [11] Macedo Maia, Siegfried Handschuh, André Freitas, Brian Davis, Ross McDermott, Manel Zarrouk, Alexandra Balahur, and Ross Mc-Dermott. Wwv’18 open challenge. pages 1941–1942, 2018.
- [12] Burton G. Malkiel. The efficient market hypothesis and its critics. *Journal of Economic Perspectives*, 17:59–82, 12 2003.
- [13] Pekka Malo, Ankur Sinha, Pekka Korhonen, Jyrki Wallenius, and Pyy Takala. Good debt or bad debt: Detecting semantic orientations in economic texts. *Journal of the Association for Information Science and Technology*, 65:782–796, 7 2013.
- [14] Stephen Merity, Nitish Shirish Keskar, and Richard Socher. Regularizing and optimizing lstm language models. *6th International Conference on Learning Representations, ICLR 2018 - Conference Track Proceedings*, 8 2017.
- [15] Matthew E. Peters, Mark Neumann, Luke Zettlemoyer, and Wen Tau Yih. Dissecting contextual word embeddings: Architecture and representation. *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, EMNLP 2018*, pages 1499–1509, 8 2018.
- [16] Aliaksei Severyn and Alessandro Moschitti. Twitter sentiment analysis with deep convolutional neural networks. *SIGIR 2015 - Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 959–962, 8 2015.
- [17] Abinash Tripathy, Ankit Agrawal, and Santanu Kumar Rath. Classification of sentiment reviews using n-gram machine learning approach. *Expert Systems with Applications*, 57:117–126, 9 2016.
- [18] Ashish Vaswani, Google Brain, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 2017.
- [19] Casey Whitelaw, Navendu Garg, and Shlomo Argamon. Using appraisal groups for sentiment analysis. *International Conference on Information and Knowledge Management, Proceedings*, pages 625–631, 2005.
- [20] Yi Yang, Mark Christopher, Siy Uy, and Allen Huang. Finbert: A pretrained language model for financial communications. 6 2020.
- [21] Lei Zhang, Shuai Wang, and Bing Liu. Deep learning for sentiment analysis: A survey. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 8:e1253, 7 2018.
- [22] Yukun Zhu, Ryan Kiros, Richard Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, and Sanja Fidler. Aligning books and movies: Towards story-like visual explanations by watching movies and reading books, 2015.